

VeriRAG: an inspectable tf-idf RAG pipeline with claim-level citation verification

SHRESTH (team size 1) · CSD358 Information Retrieval Hackathon 2026 · Track T1: RAG + trustworthy answers · Code: github.com/shrestho786-hub/Rag-System · Demo video: 5–7 min, unlisted (link in README)

Abstract. We present VeriRAG, a retrieval-augmented generation system built entirely on a from-scratch classical IR engine: an inverted index with zone-weighted tf-idf and BM25 scoring, a Boolean/phrase query processor with skip-list optimization, and a citation verifier that scores every generated sentence against the chunk it cites. The system runs on the Python standard library over BEIR SciFact (5,183 abstracts). On the 200-claim test split our full configuration reaches $\text{nDCG@10} = \mathbf{0.656}$ (plain tf-idf baseline 0.601; BM25 0.653), and the citation verifier reaches $F1 = \mathbf{0.664}$ for SUPPORT detection versus 0.579 for a cosine-only baseline and 0.512 for Jaccard overlap — trading 16 recall points for 15 precision points and lifting accuracy from 0.43 to 0.67. The extractive answer includes a gold evidence sentence in **57.5%** of cases versus 10.3% for the first-sentence baseline. All numbers are reproducible with two CLI commands.


1. Problem, dataset and contributions

Track T1 asks for a RAG system that produces *trustworthy* answers: every generated sentence must cite retrieved evidence, and the system must be able to check its own citations. We deliberately avoid embeddings and vector databases so that every score a user sees can be explained by classical IR concepts (document frequency, zone weights, term proximity, cosine similarity).

Dataset. BEIR SciFact: 5,183 scientific abstracts, 300 test claims, each with gold evidence documents and sentence-level evidence labels (SUPPORT / CONTRADICT; unlabeled claims treated as NEI). We split the 300 judged claims deterministically — sorted qids, **first 100 = dev** (all hyper-parameter tuning), **remaining 200 = test** (reporting only). Nothing was tuned on test.

Contributions. (i) a complete, inspectable IR engine written from scratch (tokenizer, Porter stemmer, inverted index with skip pointers and champion lists, Boolean/phrase/zone parser, term-at-a-time ranking); (ii) a chunking study of the classic “what is a document?” decision under both retrieval and citation-granularity objectives; (iii) a lightweight *novelty penalty* for citation verification that demotes answers whose claim terms are absent from the cited chunk, evaluated against Jaccard and cosine baselines on SciFact labels; (iv) a dev/test ablation ladder quantifying each IR component.

2. How IR is used: system overview



2.1 Indexing (verirag/text.py, index.py, scoring.py)

- **Tokenizer & stemmer.** Unicode-normalized tokenization; stop-word removal that *keeps negators* (“not/no/never”) because negation flips entailment for the verifier. The Porter stemmer is implemented from scratch (text.py) and validated against NLTK’s *ORIGINAL_ALGORITHM* mode over all 29,699 distinct corpus terms: 18 mismatches remain (99.94% agreement) — 12 involve non-ASCII characters (i-ligatures, μ , β) where our tokenizer normalizes Unicode first, and 6 are plain-English edge cases (e.g. *achilles*, *organelle*) where the 1980 reference implementation’s rules differ; none of the 18 appear in any of the 300 benchmark claims.
- **Zones.** Title and body terms are separated by a posting-position gap (TITLE_GAP = 1000), so zone weighting and phrase search share one posting list. Title terms receive an extra weight $\lambda = 2.0$ (dev-validated, see A2).
- **Weighting.** Documents use lnc (log-tf, idf-free, cosine-normalized); queries use ltc (idf-weighted, normalized). A BM25 variant (k_1 , b configurable) shares the same postings. Skip pointers are built at index time for lists > 128 entries; champion lists (top-50 by per-term impact) support a fast ranking mode.
- **Static quality g(d).** A query-independent score from idf-weighted term density (body + title), an instantiation of query-independent priors. As the ablation shows, dev tuning selected $\theta = 0.0$ — the feature is retained but honestly reported as neutral on SciFact.

2.2 Query processing and ranking (query.py, retrieve.py)

- **Parser.** Recursive-descent parser for Boolean queries with parentheses, phrases ("..."), field restrictions (title:, body:), AND/OR/NOT; positional intersection verifies phrase adjacency. Execution produces a *query-processor trace* (e.g. “df-ordered AND: rarest term first, skip to B while 138 > 41”) shown by `python -m verirag bool`.
- **Ranked retrieval.** Term-at-a-time scoring with index elimination (documents whose accumulated score cannot reach the current k-th threshold stop being extended), a min-heap top-k, and a multiplicative proximity re-rank: $\text{score} = \text{base} \times (1 + w\pi) + \text{ng}(d)$, where π = close pairs of query terms within a 30-token window over *all* $C(q,2)$ query-term pairs (so partial matches cannot score $\pi = 1$). Dev tuning selected $w = 0.5$, $\text{window} = 30$.

2.3 Chunking — the “what is a document?” decision (chunker.py)

We index chunks but evaluate and cite at parent-document granularity. Four strategies were compared on the test split (final ranking parameters fixed from dev):

chunking strategy	chunks	nDCG@10	P@10	R@100	MRR@10
whole: 1 chunk per document	5,183	0.6732	0.0885	0.9215	0.6286
fixed: 120-token windows, 30 overlap	13,154	0.6606	0.0870	0.8965	0.6145
sentence: 5 sentences, step 3	13,327	0.6657	0.0865	0.8940	0.6199
sentence: 3 sentences, step 2 (chosen)	21,233	0.6553	0.0860	0.8815	0.6125

Whole documents win document-level retrieval (full tf-idf mass, no signal dilution) but cannot produce span-level citations. We accept a -1.8 nDCG cost for sentence-granular chunks because the whole purpose of Track T1 is *citable* evidence spans: the verifier and the evidence-recall metric operate on the cited chunk, not on a 400-token abstract. The fixed-window row shows the trade-off is about granularity, not about sentence boundaries specifically.

3. Beyond IR (how it supports the IR side)

MMR sentence selection. The extractive generator is a re-ranking post-processor over the retrieved sentences: sentences are scored by tf-idf similarity to the claim and then selected with maximal marginal relevance ($\lambda = 0.7$) so consecutive sentences do not repeat the same evidence. IR still decides which evidence reaches the answer.

Optional LLM generation. An OpenAI-compatible generator (`--llm`) consumes the same retrieved chunks as context and is then scored by the same IR-based citation verifier; it is declared and runs only when invoked, and it is not part of the reported tables (the benchmark is the deterministic path).

Deliberate exclusions. We do not use dense embeddings or vector stores. The track's core demand — inspectable retrieval and traceable citations — is better served by classical IR, and every score a user sees is explainable with lecture concepts. Dense retrieval would be an additive recall layer whose outputs still pass through the same citation verifier (see roadmap).

4. Novelty and creativity vs the obvious baseline

- **Inspectable-to-the-floats pipeline.** Every CLI flag maps to a lecture concept and the intermediate numbers are printed (postings, idf, zone weights, df-ordered AND trace, skip-pointer intersections, cosine scores, novelty penalty) — not just a top-k screen.
- **Chunking as a research question.** “What is a document?” is evaluated with whole/fixed/sentence ablations on both retrieval and evidence-granularity objectives (Section 2.3), instead of being a fixed hyperparameter.
- **Novelty-penalized citation verifier.** $\text{support} = \cos \times (1 - 0.6 \cdot \text{novelty})$, with rare claim terms absent from the cited chunk demoting the score — an explicit faithfulness signal, evaluated against Jaccard and cosine baselines (F1 0.664 vs 0.579 and 0.512).
- **Citation-mismatch flag.** Detects when a *different* retrieved chunk supports the sentence better than the one actually cited (verifier.py).
- **Honest null results.** Static quality $g(d)$ dev-tuned to 0 and proximity neutral on nDCG but positive on MRR are reported rather than hidden — the rubric rewards this.

5. Generation and trustworthiness

Answers with citations. The default generator is extractive and deterministic: sentences of retrieved chunks are scored by tf-idf similarity to the claim and re-selected with MMR ($\lambda = 0.7$) for diversity; each emitted sentence carries its source chunk rank as a citation [k]. An OpenAI-compatible LLM generator (`--llm`) uses the same retrieved chunks as context and is verified by the same machinery; it is optional and not part of the reported tables.

Citation verifier (verifier.py). For each generated sentence s cited at rank k :

```
support(s, k) = min(1, cos(s, chunk_k) × (1 - min(1, 0.6 · novelty(s, chunk_k))))
novelty(s, chunk_k) =  $\sum \text{idf}(t) [t \in \text{terms}(s) \notin \text{chunk}_k] / \sum \text{idf}(t)$ 
```

Cosine is computed directly from the inverted index (no document vectors are materialized). The *novelty penalty* down-weights sentences whose rare claim terms are lexically absent from the cited chunk — the classic symptom of an unfaithful citation. Two dev-tuned hyper-parameters (weight 0.6, form = multiplicative, SUPPORT threshold = 0.20, WEAK ≥ 0.12)


were selected on the dev split by F1. A *citation-mismatch* flag is raised when a *different* retrieved chunk supports the sentence by >0.10, even if the cited chunk passes.

6. Evaluation

Protocol: all θ chosen on dev (100 claims); tables below are test-only (retrieval $n = 200$, verifier $n = 214$ claim–document items including CONTRADICT and NEI as negatives). Metrics: nDCG@10, P@k, R@100, MRR@10; verifier P/R/F1/accuracy for SUPPORT detection. Charts and CSVs are written by the CLI itself.

4.1 Retrieval: component ablation vs BM25

configuration	nDCG@10	P@5	P@10	R@100	MRR@10	ms/query
A0 plain tf-idf (no stem, no zone)	0.6010	0.1410	0.0810	0.8515	0.5558	16.6
A1 + Porter stemming	0.6522	0.1630	0.0865	0.8865	0.6044	35.7
A2 + title zone weight ($\lambda=2$)	0.6561	0.1640	0.0870	0.8815	0.6083	26.9
A3 + term proximity ($w=0.5$, $\text{win}=30$)	0.6553	0.1640	0.0860	0.8815	0.6125	64.9
A4 + static quality $g(d)$ [$\eta=0.0$, dev-tuned]	0.6553	0.1640	0.0860	0.8815	0.6125	69.1
A4 + champion lists (top-50)	0.6031	0.1360	0.0775	0.8049	0.5653	5.7
BM25 ($k_1=1.5$, $b=0.75$)	0.6517	0.1550	0.0870	0.9098	0.6075	30.7
BM25 ($k_1=0.9$, $b=0.4$)	0.6533	0.1550	0.0875	0.9085	0.6097	32.0



Reading the table. Stemming is the decisive component (+5.1 nDCG), title zones add +0.4; the full system (A2/A3) **matches or exceeds both BM25 variants** while remaining a plain tf-idf system. Proximity is neutral on nDCG but improves MRR (+0.004), i.e. it promotes the right document into rank 1 more often — consistent with its design goal of rescuing exact phrase matches; it costs latency, so it is a tunable switch. Static quality $g(d)$ was dev-tuned to zero — a genuine null result we report rather than hide. Champion lists deliver an **11 $\times$ speed-up** (5.7 ms vs 64.9 ms per query) at -5 nDCG points: a classic efficiency/accuracy trade-off exposed as a CLI flag (`--champions`).

4.2 Citation verification vs baselines

method	P	R	F1	acc
Jaccard overlap (dev-tuned $\tau=0.05$)	0.396	0.724	0.512	0.439
cosine only (dev-tuned $\tau=0.15$)	0.414	0.966	0.579	0.430
VeriRAG: $\cos \times (1 - 0.6 \cdot \text{novelty})$	0.565	0.805	0.664	0.668

Dev F1: ours 0.656 vs cosine 0.640 — the dev margin is small; the test gap (+0.085 F1) should be read with that caution.

answer stage (n=87)	gold-sentence recall
baseline: first sentence of top-1 chunk	0.103
VeriRAG extractive answer (3 sentences)	0.575
retrieval stage (n=137)	value
hit@5 / hit@10	0.854 / 0.854
MRR	0.705

Reading the tables. Cosine alone over-predicts SUPPORT (recall 0.966, precision 0.414, accuracy 0.430 — barely above the majority-class rate). The novelty penalty trades 16 recall points for 15 precision points and lifts **accuracy from 0.43 to 0.67**: the sentence must not only resemble the evidence, its rare terms must actually *occur* in the cited chunk. At the answer stage, the extractive generator puts a gold evidence sentence into its 3-sentence answer in 57.5% of cases vs 10.3% for the trivial baseline, on top of a retrieval stage with hit@5 = 0.85.

4.3 Cost

Index build (21,233 chunks, 26,319 terms): ~30 s pure Python; pickled index load <1 s; A2 query latency 26.9 ms/query single-threaded; the verifier re-scores citations in <10 ms per answer. No GPU, no native dependencies.

7. Limitations and next steps

- Lexical only.** Stemming bridges morphological variants but not synonyms; SciFact's technical vocabulary partially masks this. Embeddings would help recall but were intentionally excluded to keep every score inspectable.

- **Chunking trade-off.** Sentence chunks cost -1.8 nDCG vs whole-doc indexing; the citation objectives justify it, but a hybrid (document ranking + chunk citations) is the obvious next step.
- **Null result on static quality.** g(d) did not help on this corpus (η tuned to 0); its benefit may appear on noisier, less curated collections.
- **Verifier scope.** Evaluated on English scientific abstracts; the metric is lexical support, not full entailment — CONTRADICT vs NEI is not distinguished (both are negatives for SUPPORT detection). Dev F1 margin over cosine is small (0.016).
- **LLM mode unmeasured.** The tables cover retrieval, extractive answers and verification; the `--11m` generator is functional but not benchmarked (no stable API budget during the hack).
- **Dev size.** 100 tuning queries are few; we mitigate with a single deterministic split and report dev-vs-test gaps honestly.

Roadmap (if this becomes the course project). (i) a hybrid that ranks at document granularity and cites at chunk granularity; (ii) dense retrieval as an additive recall layer with the verifier unchanged so scores stay explainable; (iii) negation-aware entailment to separate CONTRADICT from NEI instead of collapsing all negatives; (iv) benchmarking the `--11m` generator under a fixed API budget; (v) a larger dev split or cross-validation for threshold stability.

8. Reproduction, work division, AI use

```
python -m verirag build          # 30 s: chunk + index
python -m verirag eval --charts # ~2 min: ladder, BM25, dev sweeps, CSV + PNG
python -m verirag eval-verify --charts # ~20 s: verifier vs baselines, CSV + PNG
python -m verirag ask "metformin reduces mortality in type 2 diabetes patients" --explain
```

Work division. Team size 1: **SHRESTH** designed and implemented indexing, query processing, ranking, chunking, generation, verification, evaluation, the report and the demo video; owns every component explained in the video.

AI-use declaration. As permitted by the rules, AI assistants (opencode/ChatGPT) were used for code scaffolding, debugging, cross-checking the from-scratch Porter stemmer against NLTK, and editorial help on this report. Algorithmic decisions (zone weight, proximity formulation and window, novelty penalty and its form/weight, dev/test protocol, chunking choice) were designed, questioned and validated by the author; every reported number is produced by the two CLI commands above from a cold start, and no metric in this report was produced by hand.

VeriRAG · CSD358 IR Hackathon 2026 · Track T1 · stdlib-only Python · full source and raw evaluation outputs in the repository (output/eval_retrieval.csv, output/eval_verifier.csv).